CODENEXUS INDIA 2026 / ROUND 3

Brand Rule Lab

A pre-publication review workspace for brand rules, evidence, visual assets, and human approval.

LOCAL POST-DEADLINE DRAFT | Not deployed, not an on-time submission, and not a verified final release.

The problem

Marketing teams publish across people, channels, and tools. A fact can expire, a claim can outrun its evidence, or a visual can drift from approved brand assets. Brand Rule Lab puts explainable checks and accountable human decisions in one workflow.

How it works

01 GOVERN	Owners manage workspaces and invite a team. Fact Owners review versioned facts and evidence.
02 CHECK	Creators submit copy, documents, images, packaging artwork, or short video. Rules run locally; optional AI is advisory.
03 DECIDE	Creators revise a new version. A separate Reviewer examines findings and decides.
04 TRACE	Review history, claim snapshots, drift alerts, and an audit export preserve the decision trail.

Human authority stays with the team. AI is advisory. It cannot approve a review or publish a CMS item. Deterministic rules and human decisions remain authoritative.

Project status

This local worktree is uncommitted and has not been deployed or pushed. The official Round 3 deadline ended at 10:00 IST on 13 November 2026. The original Round 3 deck and detailed rubric were not found in the available checkout or attachments.

IMPLEMENTATION MAP

Three known commitments

Mapped to the improvements described in the prior Round 2 submission; the unavailable Round 3 deck and rubric may contain other criteria.

COMMITMENT	IMPLEMENTED HERE	BOUNDARY
Human-authorised remediation and CMS	Creator corrections create preserved versions with a reason. A separate Reviewer decides. Live WordPress sends `status=draft`; local adapters cover WordPress, Shopify, and Contentful.	WordPress has mocked HTTP coverage but no real-site validation. Shopify and Contentful are mocks. No publish operation exists.
Multimodal packaging and video	Image ingestion, local OCR, SHA-256 and perceptual hashes, packaging references, and bounded video keyframe OCR are implemented.	At most 12 frames from clips up to 120 seconds. Unsampled moments, motion, logo placement, layout, safe zones, and disclaimer size are not verified.
Secure collaboration and audit	Email/password flows, invitation tokens, per-brand roles, owner transfer, backend checks, tenant audit chains, and JSON/PDF exports are implemented.	Google provider redirect reached account entry but was not authenticated; Resend delivery is untested. Password-only sessions do not persist across refresh. Independent security review and external signature remain outstanding.

Core workflow

Owner creates a workspace and invites members. Fact Owner approves versioned facts and evidence. Creator submits content; review returns findings and references. Corrections are preserved and rechecked. A separate Reviewer approves or requests changes. Approval writes claim lineage, review history, and audit events; later fact changes can flag linked claims.

Database and tenant boundary

Database startup creates base tables, applies migrations 1-14 in order, then creates indexes. Existing-v11 and fresh-database rehearsals ran repeatedly against disposable files. Historical rows with no `brand_id` remain unassigned and are not exposed as a new member tenant. SQLite with one Streamlit process supports a controlled local demo, not multi-instance scaling.

EVIDENCE AND RELEASE GATE

What was verified

Validation used disposable databases, synthetic content, local services, or mocked providers unless stated otherwise.

Automated and database checks

CHECK	RESULT
Regression suite	**349 passed, 0 failed** 6 PyMuPDF/SWIG deprecation warnings. Python `compileall` passed on 13 November 2026.
Migration rehearsal	A disposable v11 copy and a separate fresh database initialized repeatedly through migration 14. The copy preserved all 73 baseline rows across 30 tables. Integrity was `ok`; foreign-key checks were empty.
Backup and audit	Disposable backup/restore preserved synthetic records and uploaded bytes; overwrite, checksum, and traversal protections were exercised.
Synthetic evaluation	60/60 deterministic relationship checks, 10/10 synthetic asset hash checks, and 30/30 deterministic injection cases. Not a live-model accuracy benchmark.
Browser and UI	Disposable browser workflow recorded Creator submission and separate Reviewer approval with a valid three-event tenant audit chain; SQLite integrity was `ok` with no FK violations. The public landing was checked in-browser. Existing-account invitation sign-in state passes its UI regression test.

Release blockers and limits

Local secrets contain Google OIDC and Resend settings. After adding Authlib 1.9.5 through the pinned Streamlit auth extra, the browser reached the configured Google account-entry page at `http://localhost:8502/oauth2callback`; no account was selected and login was not completed. Resend delivery was not tested because no recipient was selected. No WordPress test site or credentials are configured. The Round 3 rubric is unavailable. Video checks sample frames only; OCR does not certify package layout. No usability study is claimed. The deadline elapsed and the app is not deployed.

Demonstration setup

Follow README.md and docs/ROUND3_DEMO_SCRIPT.md. Set `DATABASE_PATH` and `UPLOADS_PATH` to new disposable locations and use fictional accounts. Launch from the workspace root on port 8502 to load local OIDC secrets. Never rehearse against the configured original database.

Sources: current local worktree, prior Round 2 submission and its stated roadmap, and the official CodeNexus India event page. Round 3 judging deck and rubric were unavailable for comparison.